

— THE PROBLEM

The conversations you **can't afford to get wrong.**

A job interview. An investor pitch. Asking for a raise. Breaking bad news. The talks you replay in your head, because there is nowhere safe to practice them out loud.

So people rehearse alone, or with a friend who can't really push back. Generic AI can roleplay, but it won't challenge you like the person who actually will. And where facts matter, it **invents** the details about the person across the table, so you can't even trust it to prep you.

A job interview

An investor pitch

Asking for a raise

Breaking bad news

A GENERIC AI VOICE TOOL

CANDIDATE

Stripe's valued around \$65 billion, right?

GENERIC AI

That's right, about \$65B. Great, let's move on. ✓

Actually \$159B (Feb 2026). It has no source of truth.

— THE PRODUCT

A voice-native room to practice **any hard conversation** you can't afford to get wrong.

You pick the conversation. A voice agent plays the person across the table. Where facts matter it is grounded in cited sources and catches you when you bluff. Then it scores how you carried it.

01

Pick

Choose the conversation and, where it matters, a real target to be grilled on.



02

Spar · **voice**

A live spoken rehearsal. It probes, objects, and pushes back in real time.



03

Debrief

A scorecard of how you carried it, with the receipts where facts matter.

Grounded conversations

★ the moat

Where the other side can call your bluff on a fact. Interviews, pitches, fact-backed negotiation. Grounded in **verified, cited data**. **No other voice sim has a ground-truth layer.**

Relational conversations

range

Where there is no fact to check. Hard feedback, asking for a raise, breaking bad news, the breakup you keep putting off. The agent pushes back emotionally; the debrief coaches delivery.

— WHAT MAKES IT WORK

Three things, one principle: **ground before, score after.**

01

Voice-native

A real spoken rehearsal, not a chatbox. You talk; it listens, interrupts, and pushes back like the person across the table. Replies land in about 0.2 seconds.

Vapi · Groq Llama-70B

02

Grounded where facts matter

A cited fact pack is pulled **before** the call and injected as ground truth. Bluff, and it catches you, then shows the receipt: a link to the real filing or record.

★ [Cala.ai cited sources](#)

03

Scored afterward

A post-call debrief: a readiness grade and capability radar, plus every claim reconstructed and checked. **N of M verified**, each one linked.

GLM-5.2 debrief

The architecture is the edge. Live models invent company facts, so we never fact-check live. We pre-fetch cited truth, ground the agent, and run the rigorous check after the call.

grounding before · scoring after

— A REAL CALL · THE IB INTERVIEW

She says \$65B. The room knows it's **\$159B**, and shows the receipt.



Managing Director
 Investment Banking · live

★ grounded in 19 cited facts

rec 04:12
mute
End interview

LIVE TRANSCRIPT

MD "And how would you frame its scale for an investor?"

YOU "It's private. Last I checked, valued around **\$65 billion** after a down round."

MD "You sure about that number? Walk me through it."

► **BLUFF CAUGHT · HIGH**

DOUBLED-DOWN

Most recent valuation is **\$159B** (Feb 2026 tender). \$65B was a 2023 employee tender, not the latest.

➤ finance.yahoo.com · [Stripe valuation jumps to \\$159B](#)

readiness

C-

45 / 100

2 bluffs caught,
both sourced.

Pre-fetched cited fact pack → injected into the live agent → flagged in the post-call debrief. **No live fact-checking.** That's what keeps it fast and grounded.

WHY THIS WINS

Walk into the conversation **ready.**

Greenroom is a voice-native room to practice any hard conversation you can't afford to get wrong, grounded in cited sources where facts matter, scored when you're done. The ground-truth layer is the moat: **no other voice sim has one.**

- ✓ **Built end-to-end.** Pick → live voice spar → sourced, scored debrief. Live browser call verified.
- ✓ **Both sponsor prizes, one build.** Cala.ai is the cited ground truth; Vapi is the voice.
- ✓ **A category, not a one-off.** New conversations are a config entry behind a grounded flag, so the vision scales.

THE ONE BEAT THAT LANDS

A candidate says the "obvious" thing a generic AI would agree with. Greenroom catches it, with a dated source link.

➤ That's the moment Cala beats GPT-alone.

Practice the conversation before it counts.

Voice-native rehearsal · grounded in cited sources · scored when you're done.

★ Greenroom